

**Project Title:**

Memory Latency Profiling and Software Prefetching on the Raspberry Pi 5

---

**Project Goal:**

The goal of this project is to empirically characterize the memory hierarchy of the Raspberry Pi 5 and to evaluate the effectiveness of software prefetching as a technique to mitigate memory access latency. Using custom C++ microbenchmarks, the project will measure memory access behavior across cache levels and main memory, then analyze how explicit software prefetching impacts performance on memory-bound workloads.

---

**Hardware Requirements:** - Raspberry Pi 5 (ARM Cortex-A76 CPU) - Active cooling solution (heatsink and/or fan) to prevent thermal throttling during sustained measurements - Optional: USB power meter for optional energy/performance discussion

---

**Software Requirements:** - Operating System: Raspberry Pi OS (64-bit recommended) - Programming Language: C++ (for all benchmarks and experiments) - Tools: - `perf` (hardware performance counters for cycles, instructions, cache misses, IPC) - Python with matplotlib (for plotting and visualization of results) - Optional: `valgrind` / `cachegrind` for supplementary memory behavior analysis (non-cycle-accurate)

---

## Project Deliverables

### Deliverable 1: Memory Latency Profiler (C++ Microbenchmark)

- A custom C++ microbenchmark that:
  - Allocates memory buffers of increasing working set sizes
  - Uses a controlled access pattern (e.g., dependent accesses or pointer chasing) to expose true memory latency
  - Measures average memory access latency for each working set size
  - Output data recorded in a structured format (e.g., CSV)
- 

### Deliverable 2: Memory Hierarchy Characterization

- Plots of average memory access latency versus working set size
  - Identification of distinct latency regions corresponding to:
    - L1 cache
    - L2 cache
    - Main memory (DRAM)
  - Discussion of observed transition points between cache levels
-

### **Deliverable 3: Baseline Memory-Bound Workload**

- Selection and implementation of a memory-sensitive workload in C++ (e.g., matrix multiplication variant, array traversal kernel, or graph-style access pattern)
  - Baseline version without explicit software prefetching
- 

### **Deliverable 4: Software Prefetching Implementation**

- A modified version of the same workload that includes manual software prefetching
  - Clear documentation of:
    - Prefetch placement within the code
    - Chosen prefetch distance
    - Rationale for the chosen strategy
- 

### **Deliverable 5: Performance Evaluation and Comparison**

- Quantitative comparison between baseline and prefetch-enabled implementations using:
  - Execution time
  - Instructions per cycle (IPC)
  - Cache and memory-related performance counters from `perf`
  - Plots and tables summarizing performance differences
  - Notes on experimental controls (consistent runs, cooling, CPU frequency behavior, optional core pinning)
- 

### **Deliverable 6: Analysis and Final Report**

- Interpretation of measured results in terms of computer architecture concepts
  - Discussion of:
    - When and why software prefetching improves performance
    - Cases where prefetching provides little benefit or degrades performance
    - The relationship between memory access patterns, cache behavior, and latency hiding
    - Connection of observed behavior to underlying microarchitectural design principles
- 

### **Optional Extension: Adaptive Software Prefetching**

- Implementation of a simple adaptive prefetcher that:
  - Detects memory access stride at runtime
  - Adjusts prefetch distance dynamically
  - Comparison against static prefetching strategies
  - Extended analysis discussing trade-offs and effectiveness
-

**Project Summary:**

This project will use carefully designed C++ microbenchmarks on real ARM hardware to expose memory hierarchy behavior and evaluate software-based techniques for mitigating memory latency. The results will demonstrate how architectural constraints influence performance and how software can be structured to better exploit modern CPU designs.